

Princeton GPU Hackathon

EcoSLIM

Jun Zhang, U of Arizona
Hoang Tran, Princeton U
Chen Yang, Princeton U

Mentors

Troy Comi, Princeton U
Carl Ponder, NVIDIA

Goals

	P(:,1)	P(:,2)	P(:,3)	P(:,4)	P(:,5)	P(:,6)	P(:,7)	P(:,8)
1								1	
2								0	
3								0	
4								1	
5								1	
6								0	
7								1	
8								1	
9								0	
10								1	
....									




Figure 39-10 Scan and Scatter

```
template <class T, bool isBackward>
__global__ void compactData(T
                                size_t
                                const unsigned int
                                const unsigned int
                                const T
                                unsigned int
                                *d_out,
                                *d_numValidElements,
                                *d_indices, // Exclusive Sum-Scan Result
                                *d_isValid,
                                *d_in,
                                numElements)
{
    if (threadIdx.x == 0)
    {
        if (iGlobal < numElements && d_isValid[iGlobal] > 0) {
            d_out[d_indices[iGlobal]] = d_in[iGlobal];
        }
    }
}
```

Scatter in cudpp

1D array

- 1) For scan: writing a wrapper to call **exclusive_scan** in Thrust which is a CUDA C++ template library;
- 2) For scatter: modifying the kernel from cudpp to handle 2D particle array P;
- 3) For scatter: writing a wrapper to call this CUDA C++ kernel in Fortran; Or probably rewrite it in CUDA Fortran;
- 4) Test and profile scan and scatter using restart files generated in previous simulations;
- 5) Using cmake, Nsight, etc.

Wrapper Functions

```
file  thrust_module.cuf  cscan.cu 9+ x
chenyang > Desktop > GPU_Hackathon_Princeton > test > thrust_scan > thrust_module.cuf > ...
#include <thrust/host_vector.h>
#include <thrust/device_vector.h>
#include <thrust/scan.h>
extern "C" {
    void scan_int_wrapper( int *data_in, int N, int *data_out)
    {
        thrust::device_ptr<int> dev_ptr_in(data_in);
        thrust::device_ptr<int> dev_ptr_out(data_out);
        thrust::exclusive_scan(dev_ptr_in, dev_ptr_in+N, dev_ptr_out);
    }
    void scan_float_wrapper( float *data_in, int N, float *data_out)
    {
        thrust::device_ptr<float> dev_ptr_in(data_in);
        thrust::device_ptr<float> dev_ptr_out(data_out);
        thrust::exclusive_scan(dev_ptr_in, dev_ptr_in+N, dev_ptr_out);
    }
    void scan_double_wrapper( double *data_in, int N, double *data_out)
    {
        thrust::device_ptr<double> dev_ptr_in(data_in);
        thrust::device_ptr<double> dev_ptr_out(data_out);
        thrust::exclusive_scan(dev_ptr_in, dev_ptr_in+N, dev_ptr_out);
    }
}
```

```
file  thrust_module.cuf x  cscan.cu 9+
chenyang > Desktop > GPU_Hackathon_Princeton > test > thrust_scan > thrust_module.cuf > ...
module thrust
interface thrustscan

    subroutine scan_int(input,N,output) bind(C,name="scan_int_wrapper")
    use iso_c_binding
    integer(c_int),device:: input(*)
    integer(c_int),device:: output(*)
    integer(c_int),value:: N
    end subroutine

    subroutine scan_float(input,N,output) bind(C,name="scan_float_wrapper")
    use iso_c_binding
    real(c_float),device:: input(*)
    real(c_float),device:: output(*)
    integer(c_int),value:: N
    end subroutine

    subroutine scan_double(input,N,output) bind(C,name="scan_double_wrapper")
    use iso_c_binding
    real(c_double),device:: input(*)
    real(c_double),device:: output(*)
    integer(c_int),value:: N
    end subroutine

end interface
end module thrust
```

```
program testsort
```

```
use thrust
```

```
real, allocatable :: cpuData(:, :)
real, allocatable, device :: gpuData(:, :)
integer :: N, count
```

```
N=6
```

```
allocate(cpuData(N,2))
allocate(gpuData(N,2))
```

```
count=0
```

```
do i=1,N
  do j=1,2
    count=count+1
    cpuData(i,j)=count
  end do
end do
```

```
write(*,'(a,6(f5.1,1x))') "Before sanning", cpuData(:,1)
write(*,'(a,6(f5.1,1x))') "Before sanning", cpuData(:,2)
```

```
gpuData=cpuData
```

```
call thrustscan(gpuData(:,2),size(gpuData(:,2)),gpuData(:,2))
```

```
cpuData=gpuData
```

```
write(*,'(a,6(f5.1,1x))') "After sanning", cpuData(:,1)
write(*,'(a,6(f5.1,1x))') "After sanning", cpuData(:,2)
```

```
end program
```




Figure 39-10 Scan and Scatter

Results

Currently Loaded Modulefiles:

1) nvhpc/21.1 2) cudatoolkit/11.1 3) openmpi/

Before sanning 1.0 3.0 5.0 7.0 9.0 11.0

Column 2 Before sanning 2.0 4.0 6.0 8.0 10.0 12.0

After sanning 1.0 3.0 5.0 7.0 9.0 11.0

Column 2 After sanning 0.0 2.0 6.0 12.0 20.0 30.0

Column 2

Problems

- What problems are you currently facing?

MPI failed

```
program testsort
  use thrust
  use mpi
  use mpiDeviceUtil

  real, allocatable :: cpuData(:, :)
  real, allocatable, device :: gpuData(:, :)
  integer :: N, count, deviceID
  integer :: rank, nproc, ierr

  call mpi_init(ierr)
  call mpi_comm_size(mpi_comm_world, nproc, ierr)
  call mpi_comm_rank(mpi_comm_world, rank, ierr)

  call assignDevice(deviceID)

  N=6; count=0
  allocate(cpuData(N,2))
  allocate(gpuData(N,2))

  do i=1,N
    do i=1.2
```

← Assign GPU

Currently Loaded Modulefiles:

```
1) nvhpc/21.1  2) cudatoolkit/11.1  3) openmpi/nvhpc-21.1/4.1.0
srun: error: della-i14g2: tasks 2-3: Illegal instruction (core dumped)
srun: launch/slurm: _step_signal: Terminating StepId=34826029.0
srun: error: della-i14g1: tasks 0-1: Illegal instruction (core dumped)
```

~